

Exercise 0

Gregory Attra

02.20.2022

CS 7180 - Prof. Amato

Code Setup:

1. Unzipped the 'ex4.zip' file
2. 'cd' into the 'code' directory
3. run 'source ./init' to setup the pyenv

Questions:

1. (a)
 - Prior to the infinite loop, create a dictionary for storing the number of times state s has been reached:
 $N(s) = 0 \forall s \in S$
 - After the check for "first-visit", increment the number of times state s has been reached and use the updated $N(s)$ value to average the return at that state:
 $N(s) = N(s) + 1$
 $V(s) = V(s) + \frac{1}{N(s)}[G - V(s)]$
- (b)
 - Prior to the infinite loop, create a dictionary for storing the number of times action a has been taken at state s :
 $N(s, a) = 0 \forall s \in S \forall a \in A$
 - After the check for "first-visit", increment the number of times action a has been taken in state s and use the updated $N(s, a)$ value to average the return for that state-action pair:
 $N(s, a) = N(s, a) + 1$
 $Q(s, a) = Q(s, a) + \frac{1}{N(s, a)}[G - Q(s, a)]$
2. (a) I would expect the average return for a state to be the same. When using every-visit MC, we include returns which are computing later in the episode, in other words: smaller returns. This brings our average lower than had we used first-visit, as we would have a higher likelihood of encountering a state earlier in the episode, and thus having a higher return. However, in blackjack, episodes are incredibly brief with often fewer than five steps per episode. The probability of encountering the same state twice in a given episode is low, and even still, the brief nature of each episode means the expected return when a given state is first entered will be roughly equal, but slightly greater, than the expected return when that state is reached later in the episode.
- (b) i. First Visit:
T - 1: $G = \gamma G + r = 1$
 $Q(T, a) = Q(T, a) + \frac{1}{N(T, a)=1}[G - Q(T, a)] = 1$
T - 2 \rightarrow 9 : $G = \gamma G + r$

T - 10: $G = G + r = 9 + 1 = 10$ $Q(NT, a) = Q(NT, a) + \frac{1}{N(T,a)=1}[G - Q(NT, a)] = 10$
The estimate for $Q(NT, a) = \frac{1}{N(s,a)=1}[10 - Q(NT, a)] = 10$

ii. Every Visit:

$$Q(T, a) = 1$$


$$Q(NT, a) = \frac{G_{T-2} + G_{T-3} \dots + G_{T-10}}{N(NT, a)} = \frac{2+3+4 \dots +10}{9} = 6$$

3. Blackjack:

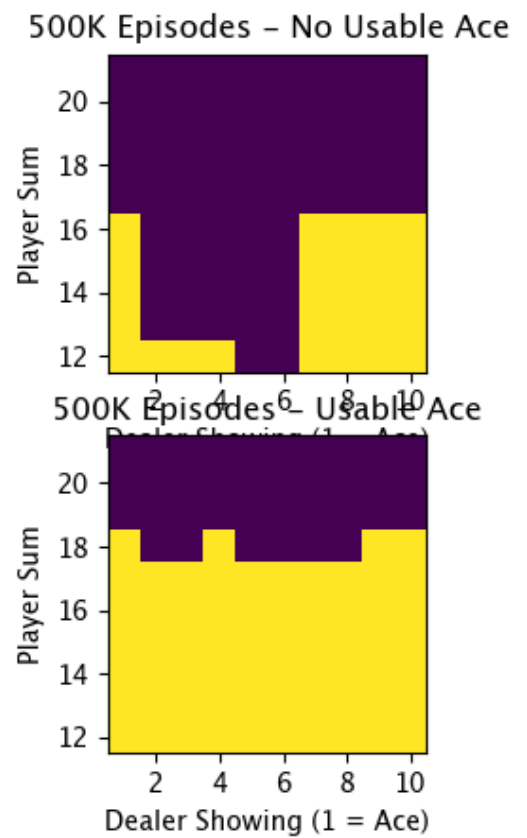
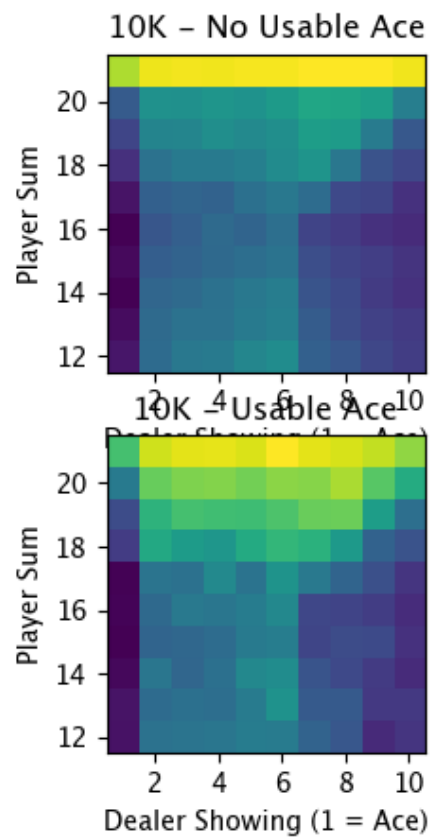
The algorithms for MC control live in ‘blackjack.py’

(a) To run Blackjack without exploring starts run: ‘python src/monte_carlo/run_blackjack.py’

(b) To run Blackjack with exploring starts run: ‘python src/monte_carlo/run_blackjack_es.py’



Blackjack First-Visit MC




Blackjack First-Visit MC w/ Exploring Starts

4. Four Rooms:

The algorithms for MC Control live in ‘algorithms.py’. The environment for the Four Rooms problem lives in ‘env.py’.

- (a) To run the four rooms problem: ‘python src/monte_carlo/run_four_rooms.py’




Four Rooms MC with Epsilon Soft Policy


5. (a)
 - (b) Because if $A_t \neq \pi(S_t)$, we exit and do not update the weight. It is therefore redundant to check $\pi(A_t|S_t)$ as we know it is 1 (since π is a greedy policy), else we would have exited.
6. Racetracks:

The algorithms for on-policy and off-policy MC Control live in ‘algorithms.py’. The environment lives in ‘racetracks.py’

 - (a) To run the racetracks problem using On-Policy MC Control: ‘python src/monte_carlo/run_on_policy_racetrack.py’
 - (b) To run the racetracks problem using Off-Policy MC Control: ‘python src/monte_carlo/run_off_policy_racetrack.py’



On-Policy MC Control for Racetrack Problem



Off-Policy MC Control for Racetrack Problem